

FLUID: A Protocol for the Agentic Data Enterprise

A Research Paper on the Federated Layered Unified Interchange Definition (v1.0)

Abstract

The enterprise is on the cusp of a paradigm shift, moving from process automation to an **Agentic Ecosystem** where autonomous AI agents drive operations. This transition, enabled by communication standards like the Model Context Protocol (MCP), exposes a foundational crisis: our data infrastructure is built on a patchwork of brittle, imperative pipelines, making it fundamentally unready for agentic consumption. This paper introduces [FLUID](#) (**Federated Layered Unified Interchange Definition**), a novel, open-source specification designed to solve this crisis. FLUID reframes data from a pipeline byproduct into a first-class **Data Product** with a clearly defined interface, contract, and implementation. By providing a declarative, version-controlled, and unified definition for every data asset, FLUID establishes the trustworthy, governable, and scalable data fabric necessary for the agentic era. This paper details the "why, what, how, when, and who" of the FLUID specification, analyzes the challenges to its adoption, and makes the case that it is the missing protocol required to mature the modern data stack.


1. The Story: From Rigs of Oil to an Ocean of Data

To understand why a protocol like FLUID is necessary, we must first understand the story of the last two decades of digital infrastructure.

The API revolution had us build magnificent "**oil rigs**" out at sea—our microservices and REST APIs. We became incredibly proficient at drilling for the valuable "oil" of transactional data, and this API-driven economy successfully fueled the app and smartphone revolution. Our enterprise oceans are now dotted with these powerful, but static, structures. Each rig is a silo, custom-built for a specific purpose, connected by a complex and often brittle network of underwater pipelines.

Th paradox of success

"The API revolution had us build magnificent 'oil rigs'—our microservices—that successfully drilled for the API 'oil' fueling the app economy, leaving our enterprise oceans dotted with these powerful, but static, structures."



But now, a fundamental shift is occurring. Data is no longer just the oil we extract; it has become **the entire ocean**. The rise of AI and the agentic era has created a rising tide of data complexity, real-time demands, and contextual awareness that our fixed rigs were never designed to handle.

We Need To Ride the Tide of Data

Embrace Agility: Build for Change

To stay afloat, our infrastructure must be as dynamic as the data that flows through it

Mandate Trustworthiness: Make Contracts Non-Negotiable


Agents cannot operate on hope. Trust must be an explicit, machine-readable contract

Engineer for Resilience

Decouple producers from consumers through well-defined data product interfaces. This isolates failure domains

Design for Scalability: Automate Governance & Enable Self-Service

The inevitable collapse of outdated data systems will occur due to the rise of AI and data demands.



This rising tide is sinking the old infrastructure. The thundering herd of agentic requests overwhelms our static APIs. The brittle pipelines break under the strain of new analytical demands. The lack of a shared language between our rigs creates data swamps and governance nightmares. To survive, enterprises need a new kind of vessel—one that is **fluid**, adaptable, and designed to ride the waves, not resist them. This is the world FLUID was conceived for.

2. What is FLUID? A Declarative Protocol for Data Products

FLUID is a declarative specification, written in YAML and managed in version control, that defines a data product in its entirety. It is not an execution engine; it is a **universal language** for an "Agentic Executor" or a compliant ecosystem of tools to understand and act upon.

Its name, **Federated Layered Unified Interchange Definition**, encapsulates its core philosophy:

- **Federated:** FLUID files are designed to be distributed, co-located with the code and teams that own the data product. This enables a true, domain-oriented data mesh architecture.
- **Layered:** The protocol explicitly supports the creation of data products in logical layers (e.g., Bronze, Silver, Gold), allowing for a clear progression from raw,

unprocessed data to refined, analytics-ready assets.

- **Unified:** It is the first standard to unify the entire data product lifecycle—interface, dependencies, build logic, contracts, and access policies—into a single, cohesive definition.
- **Interchange Definition:** It serves as a common format for exchanging data product definitions between disparate tools, teams, and agents.

The core innovation of FLUID is its **product-first structure**. A `.fluid.yml` file is organized around two primary interfaces:

1. **exposes (The Output Port):** This block defines the public interface of the data product. It is what consumers see and interact with. It contains the product's physical location, its contract (schema, quality, privacy), and its accessPolicy.
2. **consumes (The Input Port):** This block defines the data product's dependencies, whether they are physical sources (like a Kafka topic or an S3 bucket) or other FLUID data products.
3. **build (The Implementation):** This block contains the logic for *how* the output is created from the inputs. It defines the transformation engine (e.g., dbt, Spark), the execution runtime (e.g., Airflow), and stateManagement for incremental processing.

This structure cleanly separates the "**what**" (the product's interface) from the "**how**" (its internal construction), a critical distinction for creating a scalable and understandable data ecosystem.

3. Multi-Stakeholder Value: Why FLUID Matters

FLUID is designed to create value across the entire organizational chart by providing a shared language and a common framework for data.

3.1 The C-Suite Perspective: From Risk Mitigation to Innovation Velocity

For an executive, FLUID is a strategic investment in the three pillars of a modern business:

1. **Risk Reduction:** By codifying contracts and access policies, governance becomes an automated, auditable, and proactive function. This drastically reduces the risk of data quality failures impacting financial reporting and security breaches leading to costly compliance violations.
2. **Cost Efficiency:** Standardizing on a single protocol for data products eliminates redundant engineering efforts, consolidates tooling, and lowers the Total Cost of Ownership (TCO) of the data platform. It turns a chaotic collection of bespoke pipelines into an efficient "product factory."

3. **Accelerated Innovation:** The true prize is speed. When data is treated as reliable, reusable products, the time-to-market for new data-driven applications, analytics, and AI features collapses from months to days. This allows the business to react to market opportunities faster than the competition.

3.2 The Engineer's Perspective: From Brittle Pipelines to Resilient Products

For the data and analytics engineer on the front lines, FLUID solves the most persistent daily frustrations:

1. **Elimination of "Glue Code":** It replaces the thousands of lines of imperative Python and Bash scripts used to connect disparate tools with a single, declarative definition.
2. **Making dbt a First-Class Citizen:** FLUID doesn't replace dbt; it elevates it. The build block can directly reference dbt models, and the contract can inherit from a dbt schema.yml, ensuring that orchestration and transformation are always in sync.
3. **Automated Lineage & Debugging:** Because dependencies are explicit (consumes), the lineage graph is automatically generated and is always 100% accurate. Debugging becomes a structured process of validating the FLUID file's intent against the generated execution plan, rather than hunting through disconnected logs.

3.3 The Software Vendor's Perspective: The Power of a Compliant Ecosystem

For vendors across the data stack, FLUID is not a threat, but a massive opportunity. It provides a standard interface that can dramatically enhance their own value proposition.

1. **Orchestrators (Airflow, Dagster):** Instead of forcing users to manually author DAGs, a "FLUID Provider" can read a .fluid.yml file and automatically generate the corresponding DAG. This makes the platform vastly more powerful and easier to adopt.
2. **Data Catalogs (Collibra, Alation):** A FLUID connector that scans Git repositories would allow the catalog to ingest and display always-accurate lineage and metadata, solving the single biggest pain point for their customers.
3. **Data Warehouses (Snowflake, BigQuery):** Tools can be built to read the accessPolicy in a FLUID file and automatically configure the corresponding secure views and row-level security policies, making their security features more accessible.

Adopting FLUID is a go-to-market strategy that signals a product is a good citizen in the modern, composable data stack.

4. Critical Analysis: The Challenges to Adoption

A specification is only as strong as its ability to withstand scrutiny. FLUID faces three significant, non-trivial challenges.

1. **The "Agentic Executor" Implementation:** The vision depends on a compliant ecosystem of tools or a central "Agentic Executor" that can parse and act on FLUID files. Building this is a significant engineering effort. **Rebuttal:** The solution is not a single monolith, but a **decentralized, plugin-first architecture**. The FLUID standard allows a community of vendors and open-source contributors to build compliant connectors for their specific tools, fostering a vibrant ecosystem rather than a central point of failure.
2. **The "Escape Hatch" Problem:** The ability to use imperative Python scripts in the transformation block seems to undermine the declarative purity and automated lineage promises of the standard. **Rebuttal:** FLUID is pragmatic. It acknowledges that not all logic can be pure SQL. To govern this flexibility, the specification includes an optional lineage block where developers must **explicitly declare** the inputs and outputs of their scripts, making the "black box" transparent and auditable.
3. **The Cultural Inertia:** The most significant hurdle is human. Moving an organization from a pipeline-centric to a product-centric mindset requires a cultural shift, executive buy-in, and investment in training. **Rebuttal:** Adoption must be **incremental and value-driven**. FLUID's federated nature is designed for this. It can be piloted by a single, innovative team. When that team begins producing more reliable data products faster than anyone else, their success will create a gravitational pull for organic adoption across the rest of the organization.

5. Conclusion: The Foundational Layer for What's Next

The modern data stack has given us powerful, specialized tools, but it has left us with a fragmented and complex ecosystem that is ill-prepared for the agentic future. We are drowning in the complexity of our own making.

FLUID offers a path forward. By shifting the focus from imperative pipelines to declarative, version-controlled data products, it provides the structure, governance, and trust necessary for a scalable data fabric. It is the missing protocol that unifies the data stack, making it understandable, manageable, and ready for consumption by both humans and the next generation of AI agents.

The journey to a fully agentic enterprise requires a solid foundation. **FLUID is that foundation.**